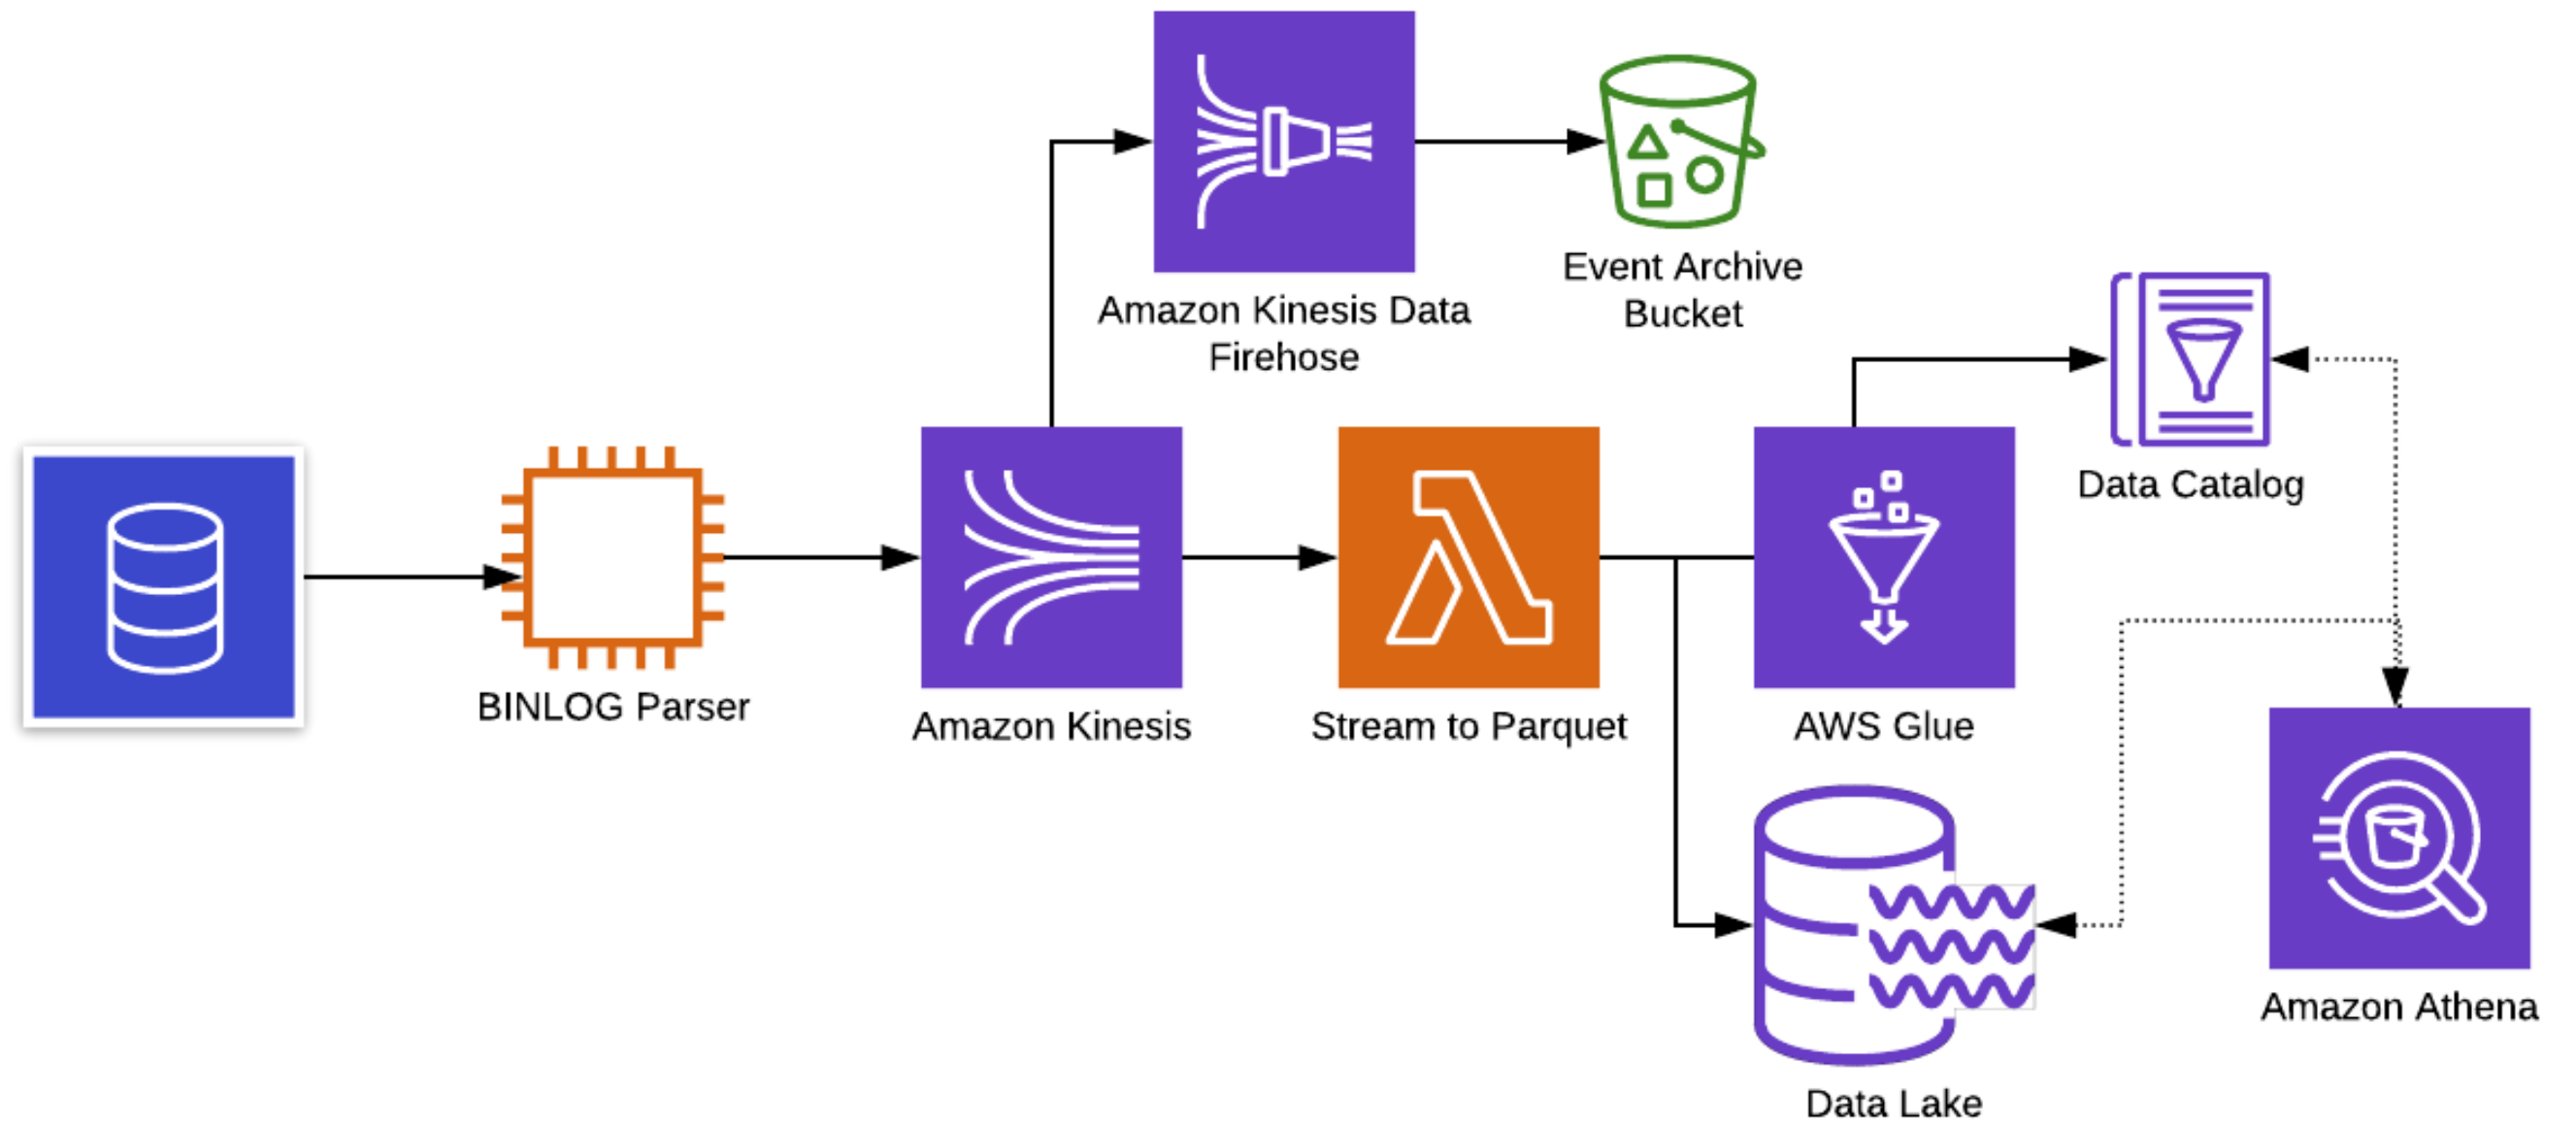



# Building History from MySQL

**IT'S BEEN A ROUGH WEEK...**



**HERE'S A BABY SLOTH IN PJ'S.**





BINLOG Parser

# **BIN LOG**

- `pip install mysql-replication`
- `GRANT REPLICATION SLAVE, REPLICATION CLIENT, SELECT  
ON *.* TO 'username'@'host'`

```
MYSQL_SETTINGS = {  
    "host": "mysql-server.some-domain.com",  
    "port": 3306,  
    "user": 'username',  
    "passwd": 'password'  
}
```

```
def main():
    # server_id is your slave identifier, it should be unique.
    # set blocking to True if you want to block and wait for the next event at
    # the end of the stream
    stream = BinLogStreamReader(
        connection_settings=MYSQL_SETTINGS,
        server_id=1337,
        blocking=True,
        only_events=[
            DeleteRowsEvent,
            WriteRowsEvent,
            UpdateRowsEvent
        ]
    )

    for binlogevent in stream:
        binlogevent.dump()

    stream.close()
```

=== TableMapEvent ===

Date: 2020-02-18T14:27:58

Log position: 28649

Event size: 65

Read bytes: 63

Table id: 91221

Schema: staging\_api

Table: activity\_feed

Columns: 11



=== WriteRowsEvent ===

Date: 2020-02-18T14:27:58

Log position: 28761

Event size: 89

Read bytes: 13

Table: staging\_api.activity\_feed

Affected columns: 11

Changed rows: 1

Values:

--

\* activity\_feed\_id : 2630701

\* loan\_id : None

\* crud : read

\* model : Loans/LoanDraft

\* primary\_key : 44206

\* user\_role : branchAdministrator

\* is\_deleted : 0

\* created\_by : 15016

\* created\_at : 2020-02-18 20:27:58.992000

\* updated\_by : 15016

\* updated\_at : 2020-02-18 20:27:58.992000

=== QueryEvent ===

Date: 2020-02-18T14:28:13

Log position: 28947

Event size: 67

Read bytes: 67

Schema: b'staging\_api'

Execution time: 0

Query: BEGIN

=== UpdateRowsEvent ===

Date: 2020-02-18T14:28:42

Log position: 30127

Event size: 293

Read bytes: 23

Table: staging\_api.inspections

Affected columns: 41

Changed rows: 1

Affected columns: 41

Values:

--

\*inspection\_id:23151=>23151

\*is\_deleted:0=>0


\*created\_by:12396=>12396

\*created\_at:2020-02-17 20:29:32.225000=>2020-02-17 20:29:32.225000

\*updated\_by:12396=>12396

\*updated\_at:2020-02-17 20:29:32.277000=>2020-02-18 20:28:42.760000

\*field\_inspector\_id:None=>None



BINLOG Parser



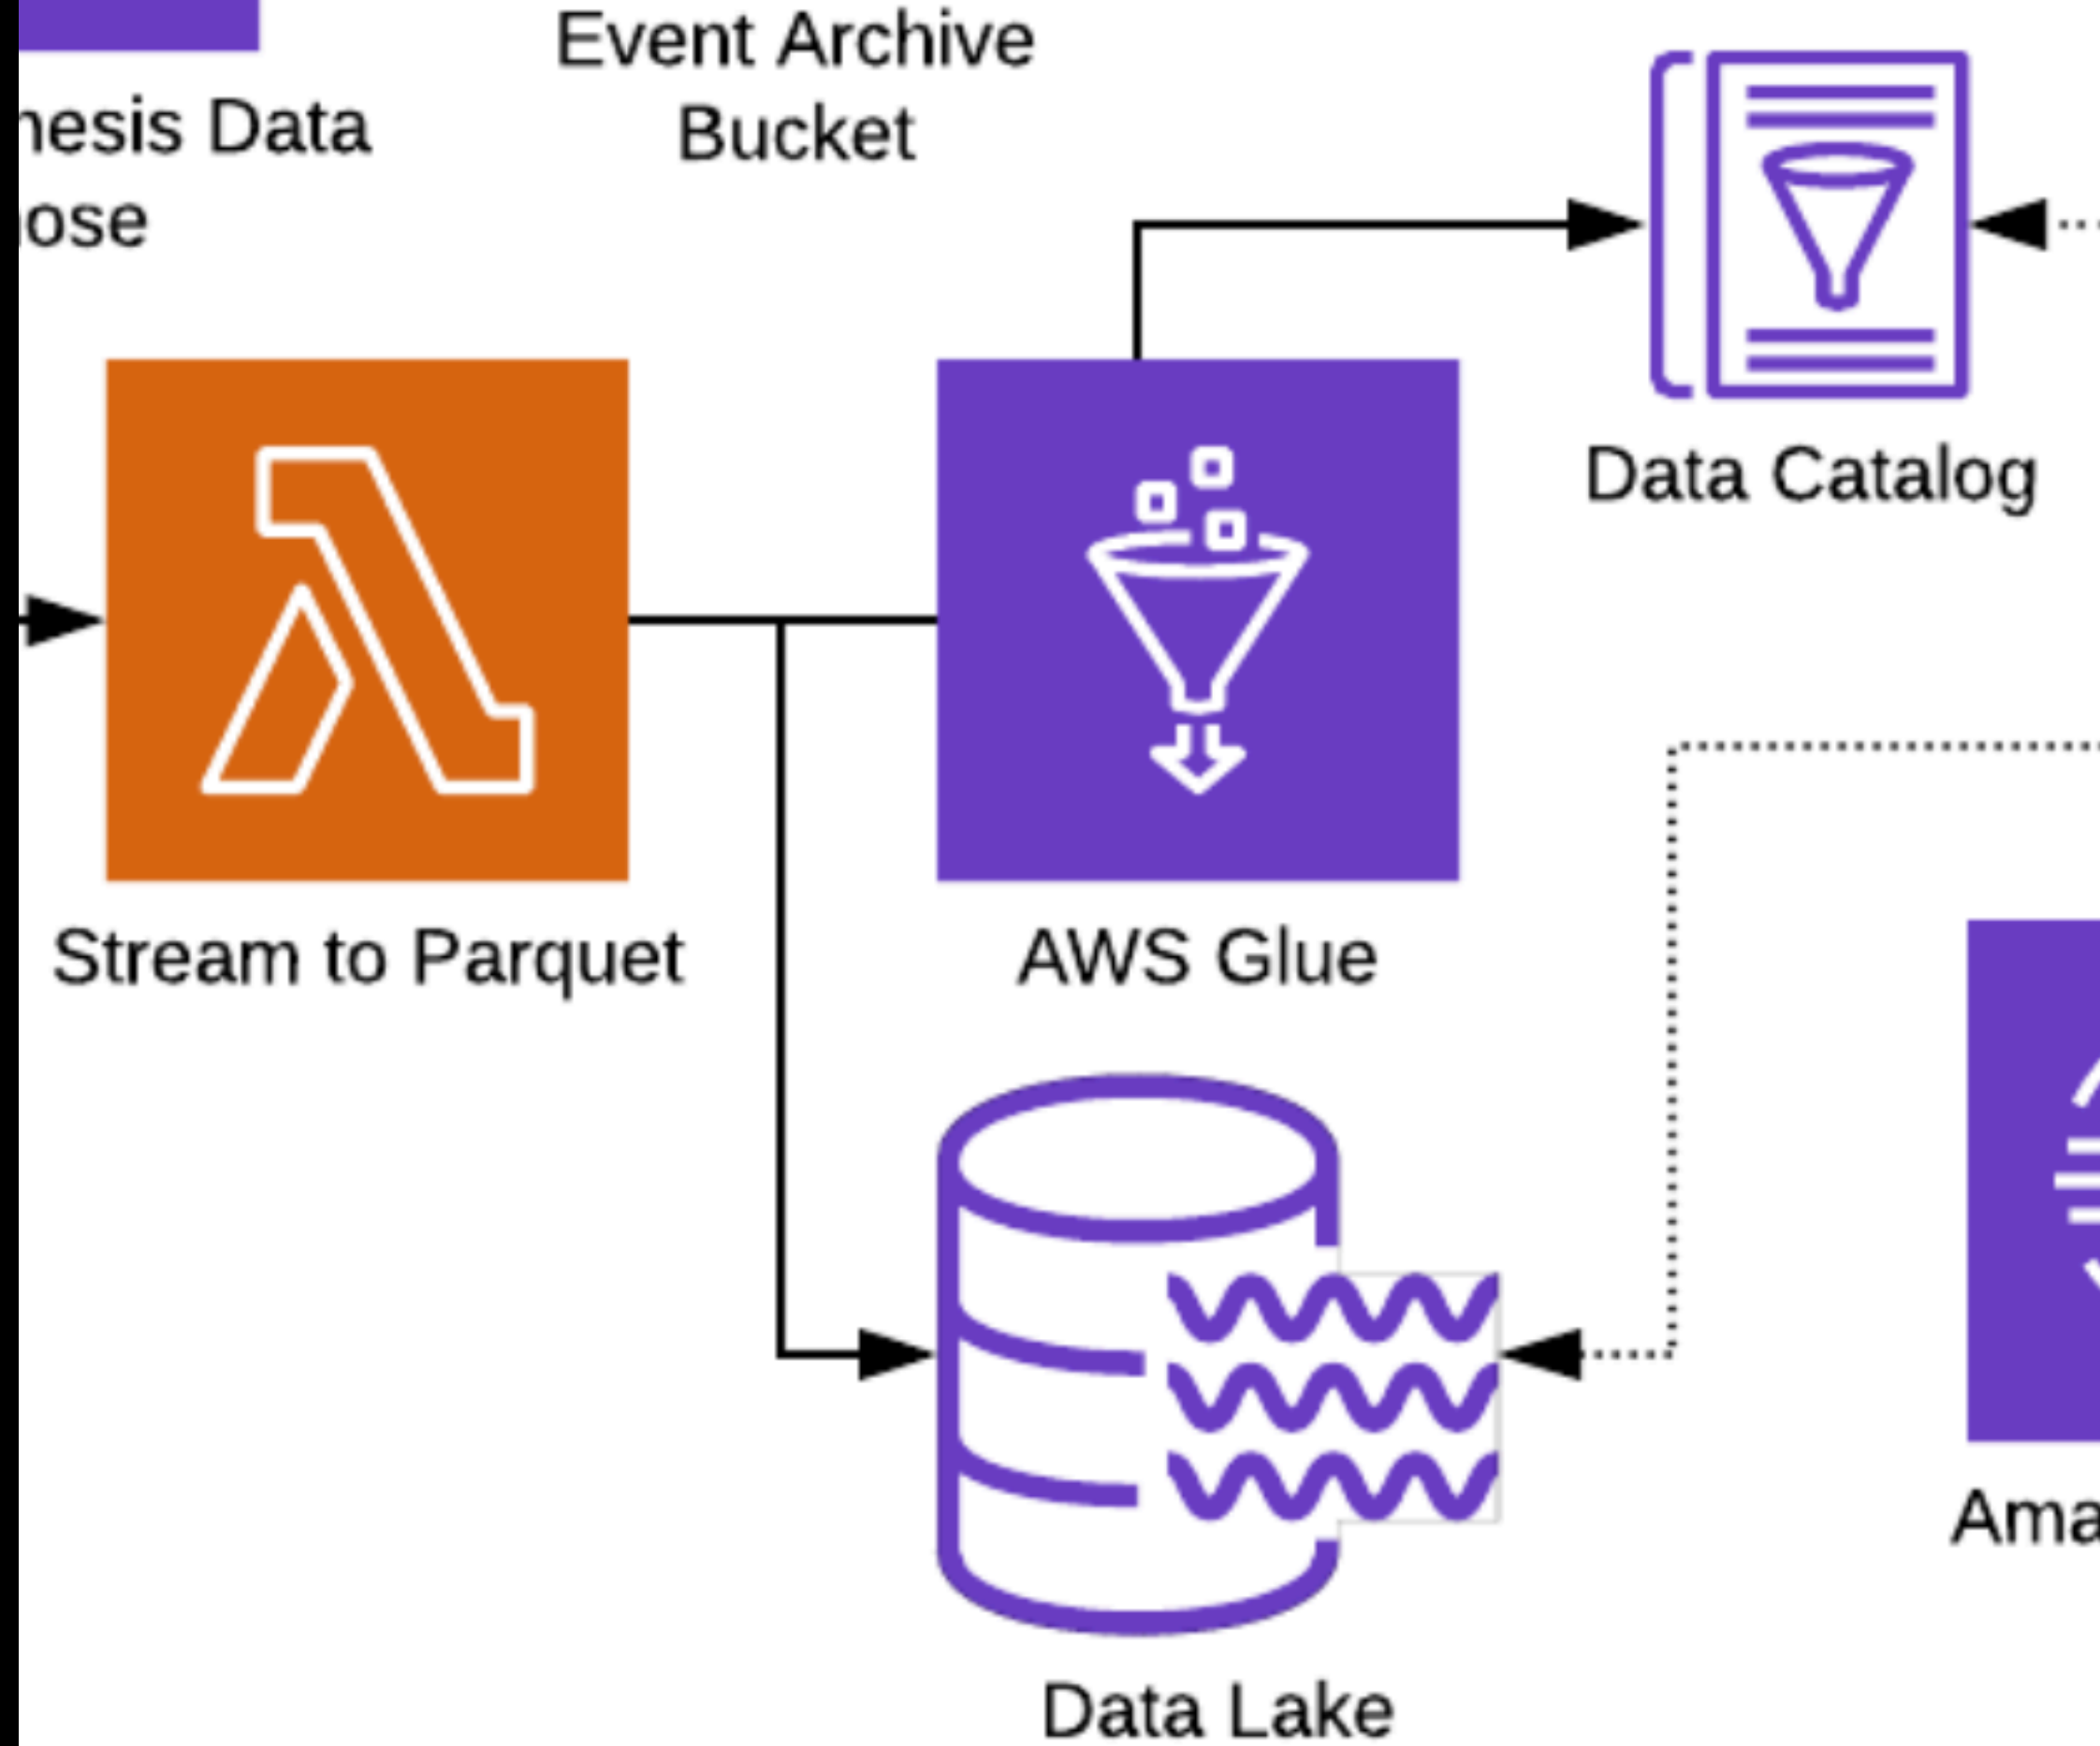
Amazon Kinesis

```
@dataclass
class ServiceEvent:
    source: str
    type: str
    data: dict
    aws_request_id: str
    version: int = field(default=1)
    event_date: str = field(default_factory=lambda: datetime.utcnow().isoformat())
    uuid: str = field(default_factory=lambda: str(uuid4()))
```

```
event_list.append(  
    ServiceEvent(  
        source=type(elem).__name__.lower(),  
        type='create',  
        data=row2dict(elem),  
        aws_request_id='test_id',  
        user_id=user_id,  
        version=2,  
    )  
)
```

```
for event in event_list:
    records_list.append(
        {
            "Data": event.to_json(),
            "PartitionKey": partition_key
        }
    )

for records_list in chunk_records(records, 500):
    try:
        kinesis_client.put_records(
            Records=records_list,
            StreamName=env.get("KINESIS_STREAM")
        )
```





# AWS Data Wrangler

- `pip install awswrangler`
- `pip install fastparquet`

```
database_name = f'data_lake'
table_name = f'{source}_{data_model}'
session = awswrangler.Session(region_name='us-east-1')
session.pandas.to_parquet(
    dataframe=event_df,
    database=database_name,
    table=table_name,
    cast_columns=cast_dict_parquet,
    path=s3_path,
    preserve_index=False,
    inplace=True,
)
```

<b>Name</b>	[REDACTED]		
<b>Description</b>			
<b>Database</b>	staging_data_lake		
<b>Classification</b>	parquet		
<b>Location</b>	<a href="#">s3://[REDACTED]api/addresses</a>		
<b>Connection</b>			
<b>Deprecated</b>	No		
<b>Last updated</b>	Fri Feb 07 14:43:05 GMT-600 2020		
<b>Input format</b>	org.apache.hadoop.hive ql.io.parquet.MapredParquetInputFormat		
<b>Output format</b>	org.apache.hadoop.hive ql.io.parquet.MapredParquetOutputFormat		
<b>Serde serialization lib</b>	org.apache.hadoop.hive ql.io.parquet.serde.ParquetHiveSerDe		
<b>Serde parameters</b>	<div>serialization.format1</div>		
<b>Table properties</b>	<div>compressionTypesnappy</div> <div>typeOfDatafile</div>		


Schema

	Column name	Data type	Partition key	Comment
1	address_id	bigint		
2	address1	string		
3	address2	string		
4	city	string		
5	state	string		
6	county	string		
7	zip	string		
8	geocoder_info	string		
9	country_code	string		
10	is_deleted	bigint		
11	created_by	bigint		
12	created_at	timestamp		
13	updated_by	bigint		
14	updated_at	timestamp		
15	uuid	string		
16	primary_key	string		
17	aws_request_id	string		
18	insights_event_type	string		
19	change_set	string		

address_id	address1	address2	city	state	county	zip	geocoder_info
8421	999 Park Place		Nantucket	GA	Troup	111	{"latitude":33.0518518,"longitude":-85.03459219999999,"location_type":"RANGE_INTERP
8421	999 Park Place		LaGrange	GA	Troup	111	{"latitude":33.0518518,"longitude":-85.03459219999999,"location_type":"RANGE_INTERP
8421	999 Park Place		LaGrange	GA	Troup	30240	{"latitude":33.0518518,"longitude":-85.03459219999999,"location_type":"RANGE_INTERP
8421	999 Park Place		Bethlehem	GA	Troup	30240	{"latitude":33.0518518,"longitude":-85.03459219999999,"location_type":"RANGE_INTERP
8421	999 Park Place		Bethlehem	GA	Barrow	30620	{"latitude":33.932589,"longitude":-83.794022,"location_type":"RANGE_INTERPOLATED





Oh yeah...  just got real



# Schema Migrations

- Default all new data to strings, and gradually fix types
- Write the schema to Kinesis (Apache AVRO, ProtoBufs)



```
def build_data_dict(tables: List, db_meta) -> Dict:
    data_dict: Dict = {"athena": {}, "python": {}}
    for table in tables:
        for column in db_meta.tables[table].c:
            if "TINYINT" in str(column.type):
                data_dict["athena"][column.name] = "bigint"
                data_dict["python"][column.name] = "int"
                continue
            if column.name == "primary_key":
                data_dict["athena"][column.name] = "string"
                data_dict["python"][column.name] = "str"
                continue
            try:
                data_dict["athena"][column.name] = python2athena(
                    column.type.python_type
                )
                data_dict["python"][column.name] = sql2python(column.type.python_type)
            except TypeError as exc_info:
                LOGGER.error(
                    "Failed to determine type for %s with message: %s",
                    column.name,
                    str(exc_info),
```

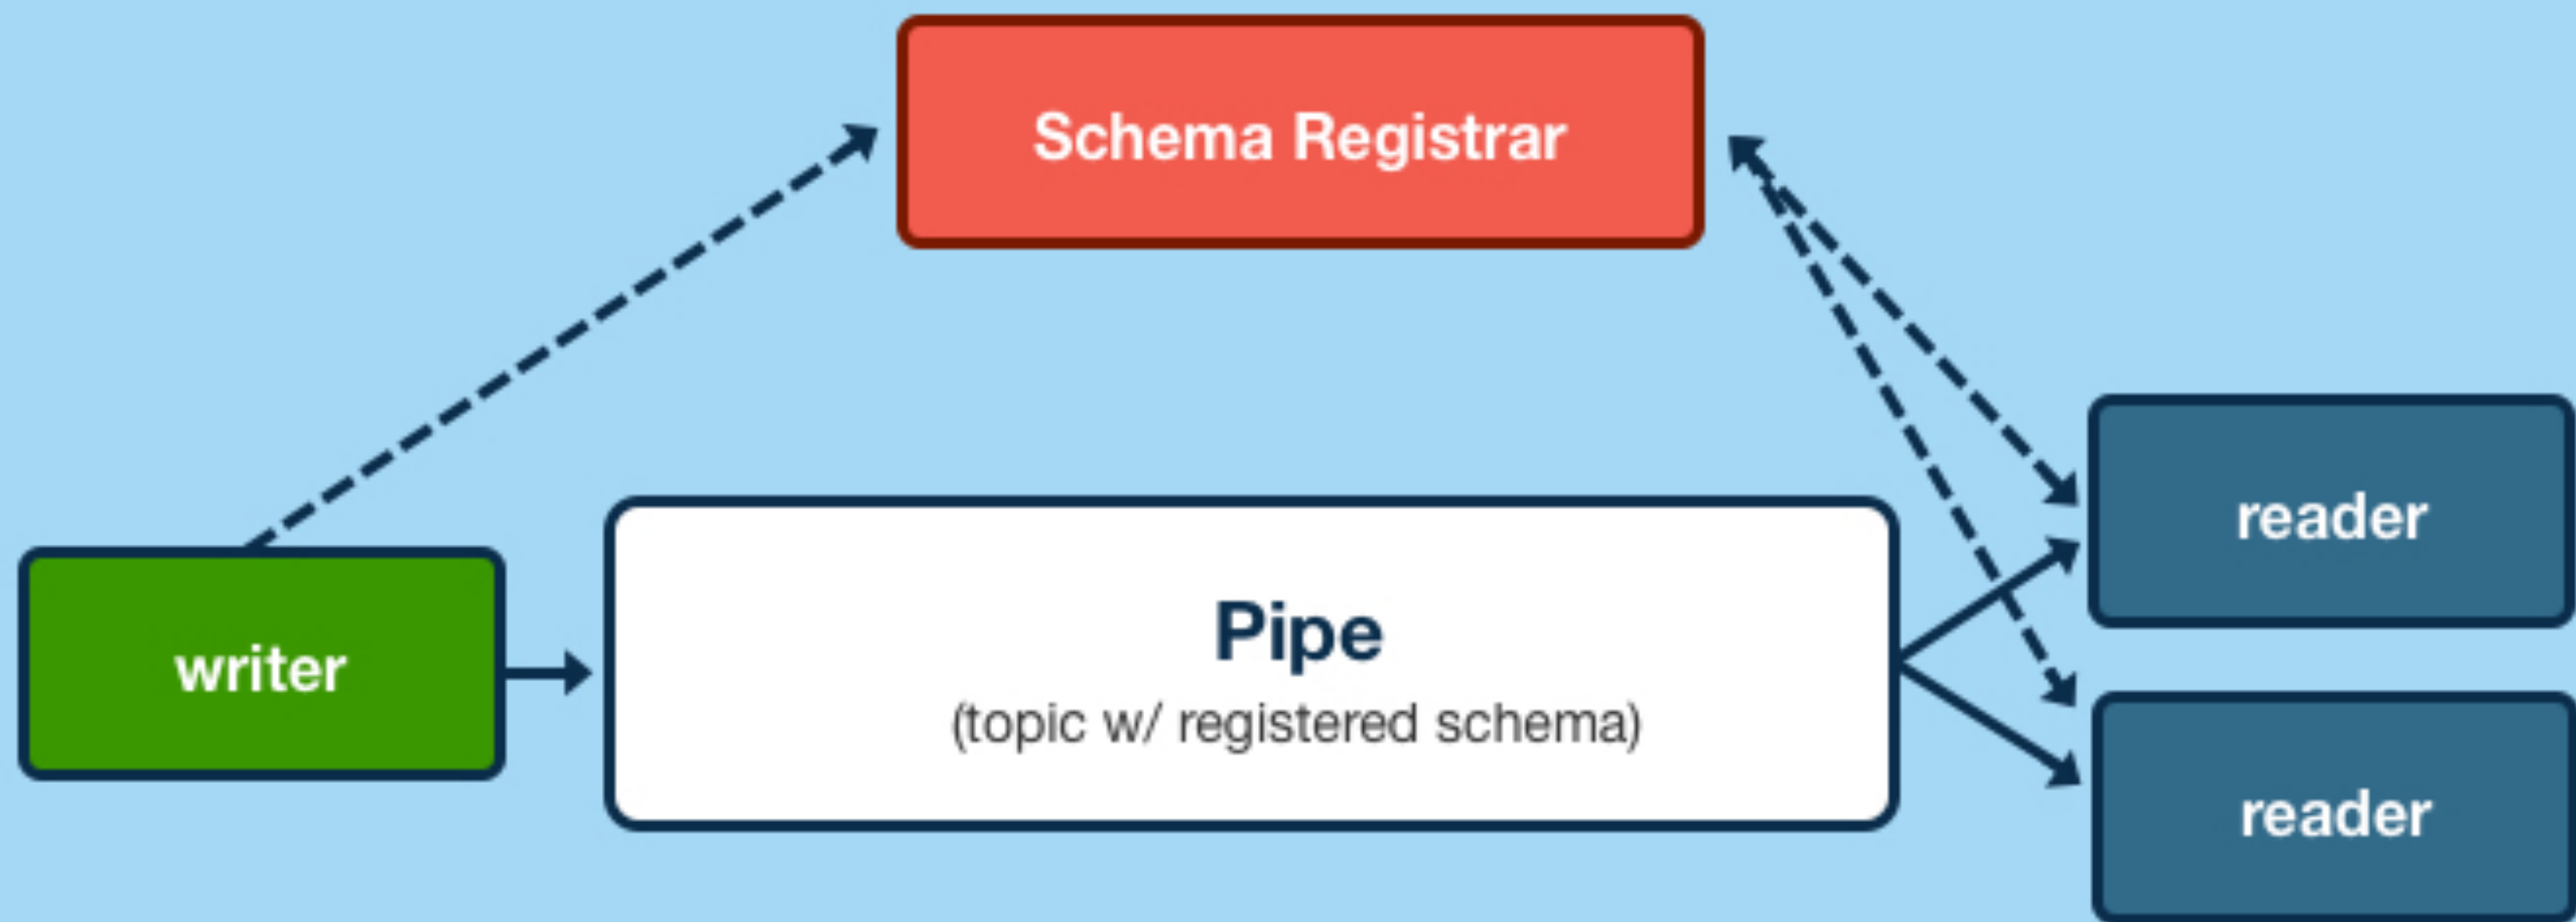
```
TYPE_MAP = {
    "<class 'int'>": {"athena": "bigint", "python": "int"},
    "<class 'float'>": {"athena": "double", "python": "float"},
    "<class 'decimal.Decimal'>": {"athena": "double", "python": "float"},
    "<class 'bool'>": {"athena": "bigint", "python": "int"},
    "<class 'str'>": {"athena": "string", "python": "str"},
    "<class 'datetime.datetime'>": {
        "athena": "timestamp",
        "python": "datetime.datetime",
    },
    "<class 'datetime.date'>": {"athena": "date", "python": "datetime.datetime"},
    "<class 'bytes'>": {"athena": "bytes", "python": "bytes"},
    "<class 'dict'>": {"athena": "string", "python": "str"},
}
```

```
def python2athena(python_type: type) -> str:
    python_type_str: str = str(python_type)
    if python_type_str in TYPE_MAP:
        return TYPE_MAP[python_type_str]["athena"]

    raise TypeError(f"Unsupported Athena type: {python_type_str}")
```

```
s3_fs = s3fs.S3FileSystem()
file_accum = []
for file_name in files:
    table = parquet.read_table(f"s3://{bucket_name}/{file_name}", filesystem=s3_fs)
    table_data_frame = table.to_pandas()
    file_accum.append(table_data_frame)
data_frame = pandas.concat(
    file_accum, keys=range(1, len(file_accum) + 1), sort=False
).reset_index(level=1, drop=True)
return data_frame
```

```
wrangler = awswrangler.Session()
group_files = wrangler.pandas.to_parquet(
    dataframe=data_frame,
    database=data_lake,
    table=f"{table_source.replace('-', '_')}_{{table_name.replace('-', '_')}}",
    preserve_index=False,
    mode='append',
    cast_columns=cast_types,
    procs_cpu_bound=1,
    procs_io_bound=1,
    inplace=True,
    path=f"s3://{bucket_name}/{table_source}/{table_name}",
)
```



- It is an error if the two schemas do not *match*.  
To match, one of the following must hold:
  - both schemas are arrays whose item types match
  - both schemas are maps whose value types match
  - both schemas are enums whose (unqualified) names match
  - both schemas are fixed whose sizes and (unqualified) names match
  - both schemas are records with the same (unqualified) name
  - either schema is a union
  - both schemas have same primitive type
  - the writer's schema may be *promoted* to the reader's as follows:
    - int is promotable to long, float, or double
    - long is promotable to float or double
    - float is promotable to double
    - string is promotable to bytes
    - bytes is promotable to string
- **if both are records:**
  - the ordering of fields may be different: fields are matched by name.
  - schemas for fields with the same name in both records are **resolved** recursively.
  - if the writer's record contains a field with a name not present in the reader's record, the writer's value for that field is ignored.
  - if the reader's record schema has a field that contains a default value, and writer's schema does not have a field with the same name, the reader's default value is used.
  - if the reader's record schema has a field with no default value, and writer's schema does not have a field with the same name, an error is signalled.
- **if both are enums:**  
if the writer's symbol is not present in the reader's enum and the reader has a `default` value, then that value is used, otherwise an error is signalled.
- **if both are arrays:**  
This **resolution** algorithm is applied recursively to the reader's and writer's array item schemas.
- **if both are maps:**  
This **resolution** algorithm is applied recursively to the reader's and writer's value schemas.
- **if both are unions:**  
The first schema in the reader's union that matches the selected writer's union schema is recursively **resolved** against it. if none match, an error is signalled.
- **if reader's is a union, but writer's is not**  
The first schema in the reader's union that matches the writer's schema is recursively **resolved** against it. If none match, an error is signalled.
- **if writer's is a union, but reader's is not**  
If the reader's schema matches the selected writer's schema, it is recursively **resolved** against it. If they do not match, an error is signalled.

## **Short Version of AVRO + JAM RULE**

- Don't rename columns
- Don't change column types
- Don't repurpose a column or a feature flag (Knight Capital Group)



**Thank you**